

Evaluation

Setup

Nineteen machines of 64 cores and 156 GiB: nine validator–committer services co-located with nine YugabyteDB nodes, three signature verifiers, one coordinator, one sidecar, one load generator. The generator embeds a mock ordering service and serves blocks to the sidecar over the Atomic Broadcast API, so these figures measure the validation and commit path with no ordering service in it. Transactions carry two read–write operations over 32 B keys, 262 B serialized, Ed25519-signed, in 10,000-transaction blocks.

A point is the highest offered rate that arrived in full, committed within 2 %, kept p99 under 1 s and held a flat in-flight count, confirmed by a 300 s hold on a fresh deployment. Queue flatness matters as much as the latency bound: a rate can be committed in full out of a growing queue, which reports the queue’s drain time as latency. Throughput counts transactions *finished*, committed plus rejected.

Committer throughput and tail latency

Throughput falls with transaction size (Fig. 1), 604,545 tps at one read–write operation to 274,364 at four: 55 % against the published 41 %. The mechanism is not the published one, which is lock contention in the coordinator’s dependency graph: here both lock-wait histograms have a count rate of exactly zero, and the graph’s cost is per key instead, about 420 ns each and flat from one key per transaction to eight.

Invalid signatures are close to free but never profitable: 558,364 tps with none, then 517,455, 519,455, 518,727 at 10 %, 20 %, 30 % — flat within 0.4 %, not the published 10 % *gain*. Driven instead at one common 559,872 tps, all three deliver it while the tail grows with the invalid share (1.1, 5.0, 7.5 s): rejection is cheap in throughput, expensive in latency.

Double spends do not reproduce, and a deployment setting is why. At the 120-way tablet pre-split this cluster uses, 0 % holds 518,000 tps and *no* rate qualifies at 5 % or above: about 20,000 tps at tens of seconds with every machine under 10 % CPU, so neither capacity nor a queue. At the default split the same shapes hold — 41,273, 30,000, 40,909 tps at 10 %, 20 %, 30 % and 334 ms to 533 ms p99, flat rather than declining — while costing the conflict-free case $1.8\times$ (518,000 against 294,893). So the split is a workload-dependent trade, not a tuning win; narrowing the graph’s chunk width refuted multi-key batching as the mechanism, leaving why 120 tablets collapse unresolved.

What latency costs at a given throughput

At 10,000-transaction blocks the median has a floor near 125 ms that does not fall with the rate, since nothing in a block moves until the block is cut (Fig. 2). At 500 transactions a block the floor is 40 ms and 380,100 tps holds at 125 ms median, 208 ms p99 — the median large blocks give at 10,000 tps, at thirty-eight times the rate. The top of that curve measures the generator, whose block preparation caps it near 850 blocks per second (425,000 tps here; 450,000 offered delivered 432,309), so the committer’s ceiling at this

size is not yet known.

End to end, with a real ordering service

TODO — running. A second arm puts a real Arma ordering service in the path (four parties, two shards, twenty further machines, committer tier unchanged). The published work has no end-to-end figure: ordering alone reaches 414,000 tps at this topology and the committer alone 419,000–474,000, which bracket it. A size sweep against its Figure 7b is queued at 300 B, 512 B, 1,024 B, 2,048 B and 4,096 B, where ordering alone reaches 413,000 down to 27,000 tps — each within 10 % of 120 MB s^{-1} , so past 300 B it is bandwidth-bound and the rate is bytes per second over size.

Threats to validity

The load generator is the instrument and twice became the limit: ECDSA signing, resolved with Ed25519, and block preparation, which caps the small-block curve. A knee is resolved to the search’s 8 % step, so it is a lower bound, drawn one-sided in Fig. 1; and the baseline drifts 10 % between days, so only same-day comparisons are used.

Committer throughput and tail latency, at a one second latency bound

Each bar is the highest rate held for 300 s with 99th percentile latency under one second, no queue growth, and the offered rate arriving. Throughput counts committed plus rejected transactions, and the hatched part of a bar is the rejected share, labelled where it is non-zero, so the two experiments are compared at a matched reject rate. The whisker is the search's 8% step, one-sided because a knee is a lower bound. The paper's series is its Figure 9: the ends as quoted in its text, the intermediate points read off its plots.




Figure 1: Throughput and 99th-percentile latency against transaction size, invalid-signature share and double-spend share, against the published figures. Hatched bars are the rejected share, so the two experiments are compared at matched reject rates rather than matched labels.




Figure 2: Latency against throughput at two block sizes; axis clipped at 400 ms, points above it labelled.